

COMP0037

Report

Planning in Uncertain Worlds

Group AS

<u>Student Name</u>	<u>Student number</u>
Arundathi Shaji Shanthini	16018351
Dmitry Leyko	16021440
Tharmetharan Balendran	17011729


Department: Department of Electronic and Electrical Engineering
Submission Date: 28th of April 2020

Contents

1	Decision Re-Plan Policy	2
1.1	Policy Selection when Obstacle is Observed	2
1.2	Policy Selection at Start	4
1.3	Considering the Probability of the Obstacle Being Present	5
1.4	Considering Multiple Obstacles	6
2	Implementation of the System in ROS	8
2.1	Planning path via an Aisle	8
	Appendices	11

1 Decision Re-Plan Policy

1.1 Policy Selection when Obstacle is Observed



(a) The original planned path from I to G (b) An obstacle in aisle B obstructs the going through Aisle B and C. planned path of the robot.

Figure 1: Illustration of case where robot observes an obstruction to its planned path.

The scenario that we will be analysing is the case shown in Fig. 1a. The robot is required to go from a cell I to a cell G . These cells are marked blue and green in Fig. 1a respectively. The figure also shows the original planned path that the robot computed going down aisle B. However, once the robot turns into aisle B it observes that the aisle is blocked. This observation is done at the point when the robot reaches the cell labelled B_1 . At this point the robot can either decide to wait until the obstruction clears or it can re-plan a path. Once the robot observes the obstacle, the time the robot must wait for the obstacle to clear may be represented by the expression in Eq. (1).

$$T = \frac{0.5}{\lambda_B} + \tilde{T} \quad (1)$$

The wait time is dependent on λ_B and a random variable \tilde{T} . The random variable \tilde{T} is sampled from a exponential distribution with a rate parameter of $2\lambda_B$. The probability density function (PDF) for \tilde{T} is shown in Eq. (2).

$$f(t) = \begin{cases} \lambda e^{-\lambda t} & t \geq 0 \\ 0 & t < 0 \end{cases} \quad (2)$$

where, the rate parameter $\lambda = 2\lambda_B$.

As previously mentioned, the robot has two options to choose from: to wait for the obstacle to clear, or to re-plan and execute the new path. The two are different policies the robot must choose from. We use the symbol π to denote a policy. A policy is a mapping from the world state to an action the robot can execute.

Let us assume that if the robot decides to wait the total path length (number of cells) will be K_1 while if the robot decides to re-plan and execute the total path length will be K_2 . We let the quantity K equal to the larger value between K_1 and K_2 . Now we may write the policy for the robot to wait as π_K^1 and the policy for re-planning as π_K^2 . These policies are padded correspondingly to produce actions \mathbf{u}_K^1 and \mathbf{u}_K^1 that are padded with zero-cost state preserving actions.

To see which policy is better on average, we consider the expected value of the cost function for both cases. The case when policy π_K^1 is chosen is characterized by the inequality shown in Eq. (3).

$$\mathbb{E} [L (\pi_K^1)] \leq \mathbb{E} [L (\pi_K^2)] \quad (3)$$




Figure 2: The path for the re-plan policy π_K^2 which bypasses aisle B and goes down aisle C.

We can see from Fig. 1a that the cost of the original planned path is given by the expression in Eq. (4). In this equation, the terms L_{XY} denote the cost of the shortest path between cell X and cell Y . Additionally, the term L_W represents the non-zero cost of the action u_w . This action is state preserving (i.e. $\mathbf{x}_{k+1} = f(\mathbf{x}_k, \mathbf{u}_w) = \mathbf{x}_k$) and is used to represent the action of the robot waiting for one unit of time.

$$L(\pi_K^1) = L_{IB_1} + TL_W + L_{B_1B} + L_{BC} + L_{CG} \quad (4)$$

From Fig. 2 which shows the re-planned path, we can also see that the cost of this path is equal to the expression in Eq. (5)

$$L(\pi_K^2) = L_{IB_1} + L_{B_1C_1} + L_{C_1C} + L_{CG} \quad (5)$$

Substituting the expressions in Eq. (4) and Eq. (5) into Eq. (3). We obtain the inequality shown in Eq. (6)

$$\begin{aligned} \mathbb{E}[L_{IB_1} + TL_W + L_{B_1B} + L_{BC} + L_{CG}] &\leq \mathbb{E}[L_{IB_1} + L_{B_1C_1} + L_{C_1C} + L_{CG}] \\ L_{IB_1} + \mathbb{E}[T]L_W + L_{B_1B} + L_{BC} + L_{CG} &\leq L_{IB_1} + L_{B_1C_1} + L_{C_1C} + L_{CG} \\ \mathbb{E}[T] &\leq \frac{L_{B_1C_1} + L_{C_1C} - L_{B_1B} - L_{BC}}{L_W} \end{aligned} \quad (6)$$


The quantity $\mathbb{E}[T]$ is the expected value for the time the robot has to wait for the obstacle to clear. As we know the distribution that the variable is sampled from we can compute the expected value. The expected value for the time taken is given by the expression found in Eq. (7)

$$\begin{aligned} \mathbb{E}[T] &= \mathbb{E}\left[\frac{0.5}{\lambda_B} + \tilde{T}\right] \\ &= \frac{0.5}{\lambda_B} + \mathbb{E}[\tilde{T}] \\ &= \frac{0.5}{\lambda_B} + \int_0^\infty 2\lambda_B t e^{-2\lambda_B t} dt \\ &= \frac{0.5}{\lambda_B} + 2\lambda_B \left[(t) \left(-\frac{1}{2\lambda_B} e^{-2\lambda_B t} \right) + \frac{1}{2\lambda_B} \int_0^\infty e^{-2\lambda_B t} dt \right]_0^\infty \\ &= \frac{0.5}{\lambda_B} + 2\lambda_B \left[(t) \left(-\frac{1}{2\lambda_B} e^{-2\lambda_B t} \right) - \left(\frac{1}{2\lambda_B} \right)^2 e^{-2\lambda_B t} \right]_0^\infty \\ &= \frac{0.5}{\lambda_B} - \left[t e^{-2\lambda_B t} + \frac{1}{2\lambda_B} e^{-2\lambda_B t} \right]_0^\infty \\ &= \frac{0.5}{\lambda_B} - (0 + 0 - 0 - \frac{1}{2\lambda_B}) = \frac{0.5}{\lambda_B} + \frac{1}{2\lambda_B} = \frac{1}{\lambda_B} \end{aligned} \quad (7)$$

Substituting the expression from Eq. (7) into Eq. (6) we obtain the inequality in Eq. (8). The right-hand side of the inequality in Eq. (8) represents the smallest possible value for λ_B for which the waiting policy π_K^1 is a better option than the re-plan policy π_K^2 . The inequality also takes into consideration the constraint that $\lambda_B > 0$ and negates the solution when $\lambda_B < 0$.

$$\begin{aligned} \frac{1}{\lambda_B} &\leq \frac{L_{B_1C_1} + L_{C_1C} - L_{B_1B} - L_{BC}}{L_W} \\ \lambda_B &\geq \frac{L_W}{L_{B_1C_1} + L_{C_1C} - L_{B_1B} - L_{BC}} \end{aligned} \quad (8)$$

1.2 Policy Selection at Start



(a) The scenario where the robot decides to go down Aisle B, encounters an obstacle and waits for it to clear. (b) The scenario where the robot decides to avoid Aisle B completely due to the obstacle.

Figure 3: The different policies the robot can pick from at the beginning.

Instead of reacting to the obstacle as the robot observes it, the robot may also make a decision before it starts to move. The robot has knowledge of where the obstacles will be and the probability distribution for the wait time. Depending on the probabilities involved, the robot may decide to avoid the obstacle altogether instead of going the route with the obstacle and wait. In the case of the warehouse example with 5 aisles, and one obstacle in aisle B, the two policies are to choose to travel down aisle B and wait if an obstacle is present (illustrated in Fig. 3a) or to avoid aisle B and plan directly along aisle C (illustrated in Fig. 3b). By avoiding travelling down aisle B before replanning the robot avoids traversing the extra distance to the cell marked B_1 which makes this approach favorable.

Similar to the approach in §1.1, we may denote the policy that goes down aisle B as π_K^1 and the policy that goes down aisle C as π_K^2 . We also consider the policy associated with choosing to go down aisle B and then replanning down aisle C as soon as the obstacle is observed. This policy will be represented as π_K^3 . Once again these are padded with zero-cost state preserving actions so as to be of the same dimension.

Fig. 3 depicts the two possible paths that represent the two former policies. The three different policies and their costs are as follows:

- Policy π_K^1 : Drive down aisle B and wait for the obstacle to clear. The cost of this scenario is given by the expression in Eq. (9).

$$L(\pi_K^1) = L_{IB_1} + TL_W + L_{B_1B} + L_{BC} + L_{CG} \quad (9)$$

- Policy π_K^2 : Completely avoid aisle B and traverse directly to the goal via aisle C. It is known that no obstacle exists on this aisle. The cost for this policy is given in Eq. (10).

$$L(\pi_K^2) = L_{IC} + L_{CG} \quad (10)$$

- Policy π_K^3 : Choose to go down aisle B and then replan down aisle C as soon as the obstacle is observed. The cost of this policy is given in Eq. (11)

$$L(\pi_K^3) = L_{IB_1} + L_{B_1C} + L_{CG} \quad (11)$$

Comparing the costs for policies π_K^2 and π_K^3 (given by Eq. (10) and Eq. (11) respectively), we can see that they only differ by only a few terms. Whereas π_K^2 goes directly down aisle C, π_K^3 first attempts to go down aisle B. This results in π_K^2 having the cost L_{IC} to get to aisle C while π_K^3 has the cost $L_{IB_1} + L_{B_1C}$ to get to aisle C. As L_{IC} is the cost of the best path, $L_{IC} \leq L_{IB_1} + L_{B_1C}$. Therefore we may ignore policy π_K^3 as the robot will always pick policy π_K^2 over policy π_K^3 .

As we are concerned with the average case scenario we compare the expected value of the loss function just as we did in section 1.1. Policy π_K^2 will be chosen over policy π_K^1 if the expected value of it's cost is lower as shown by Eq. (12).

$$\mathbb{E}[L(\pi_K^2)] \leq \mathbb{E}[L(\pi_K^1)] \quad (12)$$

Then, substituting the loss functions from Eq. (9) and (10) into the inequality in Eq. (12), we obtain the result shown in Eq. (13).

$$\begin{aligned} \mathbb{E}[L_{IC} + L_{CG}] &\leq \mathbb{E}[L_{IB_1} + TL_W + L_{B_1B} + L_{BC} + L_{CG}] \\ \mathbb{E}[T]L_W &\geq L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC} \\ \mathbb{E}[T] &\geq \frac{L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC}}{L_W} \end{aligned} \quad (13)$$

By substituting the expected value for T found in Eq. (7) we obtain a constraint for λ_B as shown in Eq. (14). The right hand side of the final inequality in Eq. (14) represents the largest value for λ_B for which the robot will decide to go directly down aisle C avoiding aisle B. Once again the sultions for $\lambda_B < 0$ have been ignored.

$$\begin{aligned} \frac{1}{\lambda_B} &\geq \frac{L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC}}{L_W} \\ \lambda_B &\leq \frac{L_W}{L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC}} \end{aligned} \quad (14)$$

1.3 Considering the Probability of the Obstacle Being Present

In reality, the obstacle would not be present there all the time. This can be taken into account using a probability, say p_B associated with the scenario that the obstacle is present¹. The mean wait time in the case the obstacle is present has been derived in Eq. (7). In the case the obstacle isn't present, the robot does not have wait and therefore the mean wait time is 0. By taking a sum of these wait times weighted by their probabilities of occurring we obtain the expected wait time that takes into consideration the probability of the obstacle being present as shown below in Eq. (15):

$$\begin{aligned} \mathbb{E}_B[T] &= p_B \cdot \mathbb{E}[T] + (1 - p_B) \cdot (0) \\ &= p_B \cdot \frac{1}{\lambda_B} = \frac{p_B}{\lambda_B} \end{aligned} \quad (15)$$

where, $\mathbb{E}[T]$ is expected time to wait given the obstacle in aisle B exist and it is given by Eq. 6

¹This means that the probability that the obstacle is absent can be given by $(1 - p_B)$

Given Eq. (15) we may reconsider Eq. (13) and substitute this new value for the expected wait time:

$$\begin{aligned}\mathbb{E}_B[T] &\geq \frac{L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC}}{L_W} \\ \frac{p_B}{\lambda_B} &\geq \frac{L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC}}{L_W}\end{aligned}\tag{16}$$

Assuming that the rate parameter λ_B is constant, we may arrive at the equation shown in Eq. (17).

$$p_B \geq \frac{\lambda_B (L_{IC} - L_{IB_1} - L_{B_1B} - L_{BC})}{L_W}\tag{17}$$

It should be noted however that this was the inequality to choose the policy to plan straight through aisle C. Therefore the R.H.S of the inequality in Eq. (17) represents the value for p_B below which the robot will attempt to drive down aisle B first.

1.4 Considering Multiple Obstacles

Let us now assume that both aisles B and C have obstacles on them. The wait time for the obstacle in aisle B to clear is sampled from the distribution outlined in Eq. (1). The probability of the obstacle in aisle B being present is once again given as p_B . Similarly the obstacle in aisle C has an identical wait time distribution as the obstacle in aisle B but with a characterizing parameter of λ_C . The probability of this obstacle being present is given by the probability p_C . Following this, the expected wait times for obstacle B and obstacle C are as follows:

$$\begin{aligned}\mathbb{E}_B[T] &= \frac{p_B}{\lambda_B} \\ \mathbb{E}_C[T] &= \frac{p_C}{\lambda_C}\end{aligned}\tag{18}$$

The robot has the same start and goal cells as described in §1.1. Given that the robot chooses a policy at the beginning, there are 6 different key scenarios that are worth considering:

- Policy π_K^1 : The robot goes down aisle B and observes an obstacle. It then decides to wait for the obstacle to clear. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^1)] = L_{IB_1} + \mathbb{E}_B [T] L_W + L_{B_1B} + L_{BC} + L_{CG}\tag{19}$$

- Policy π_K^2 : The robot goes down aisle B and observes an obstacle so re-plans down aisle C. While travelling down aisle C it observes another obstacle and waits for it to clear before continuing to the goal. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^2)] = L_{IB_1} + L_{B_1C_1} + \mathbb{E}_C [T] L_W + L_{C_1C} + L_{CG}\tag{20}$$

- Policy π_K^3 : The robot goes down aisle B and observes an obstacle so re-plans down aisle C. While travelling down aisle C it observes another obstacles so it re-plans down aisle D which is known to have no obstacles. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^3)] = L (\pi_K^3) = L_{IB_1} + L_{B_1C_1} + L_{C_1} + L_{DG}\tag{21}$$

- Policy π_K^4 : The robot goes down aisle C and observes an obstacle so it waits for the obstacle to clear before continuing to the goal. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^4)] = L_{IC_1} + \mathbb{E}_C [T] L_W + L_{C_1C} + L_{CG}\tag{22}$$

- Policy π_K^5 : The robot goes down aisle C and observes an obstacle so it re-plans down aisle D which is known to have no obstacles. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^5)] = L (\pi_K^5) = L_{IC_1} + L_{C_1D} + L_{DG} \quad (23)$$

- Policy π_K^6 : The robot directly goes down aisle D which is known to have no obstacles and avoids both aisles B and C. The expected cost for this policy is given by:

$$\mathbb{E} [L (\pi_K^6)] = L (\pi_K^6) = L_{ID} + L_{DG} \quad (24)$$

We ignore the case of going down aisle A and aisle E as these will inherently have a larger path length due to the topology of the example scenario. Additionally back-tracking of aisles is also ignored (e.g. going from aisle C back to aisle B) as this would also have larger costs than those considered above.

We are given that the robot chooses policy π_K^6 by going straight down aisle D. This suggests that the expected cost of this policy is lower than any other policy available to the robot as suggested by Eq. (25)

$$\mathbb{E} [L (\pi_K^6)] = \min_{n=1,\dots,6} \mathbb{E} [L (\pi_K^n)] \quad (25)$$

We can see that some of these policies have common terms. For example we may have a look at policies π_K^2 and π_K^4 . π_K^4 will be picked over π_K^2 if the following inequality holds:

$$\begin{aligned} \mathbb{E} [L (\pi_K^4)] &\leq \mathbb{E} [L (\pi_K^2)] \\ L_{IC_1} + \mathbb{E}_C [T] L_W + L_{C_1C} + L_{CG} &\leq L_{IB_1} + L_{B_1C_1} + \mathbb{E}_C [T] L_W + L_{C_1C} + L_{CG} \\ L_{IC_1} &\leq L_{IB_1} + L_{B_1C_1} \end{aligned} \quad (26)$$

As we know that L_{IC_1} represents the minimum path cost from I to C_1 , the inequality in Eq. (26) is true. This results in the robot always picking π_K^4 over π_K^2 . We can therefore ignore policy π_K^2 . Similarly the robot will always pick π_K^6 over π_K^3 and π_K^5 . Therefore we can ignore both policies π_K^3 and π_K^5 . We have therefore simplified the inequality in Eq. (25) as follows:

$$\mathbb{E} [L (\pi_K^6)] = \min_{n=1,4,6} \mathbb{E} [L (\pi_K^n)] \quad (27)$$

We start by comparing the policies π_K^1 and π_K^6 . If π_K^6 is chosen between the two, then the following inequality must hold:

$$L (\pi_K^6) \leq L_{IB_1} + \mathbb{E}_B [T] L_W + L_{B_1B} + L_{BC} + L_{CG} \quad (28)$$

Similarly to choose π_K^6 over π_K^4 the inequality in Eq. (29) must hold:

$$L (\pi_K^6) \leq L_{IC_1} + \mathbb{E}_C [T] L_W + L_{C_1C} + L_{CG} \quad (29)$$

The maximum value the $L (\pi_K^6)$ for policy π_K^6 to be chosen is either that given by Eq. (28) or by Eq. (29) whichever one is smaller. The inequalities cannot be merged to find the exact maximum value for the maximum path cost. However, by using logic, we may derive an upper bound for the path cost of π_K^6 . From Eq. (28) and Eq. (29) we can obtain the following inequality which will also hold:

$$\begin{aligned} L (\pi_K^6) &\leq L_{IB_1} + \mathbb{E}_B [T] L_W + L_{B_1B} + L_{BC} + L_{CG} + L_{IC_1} + \mathbb{E}_C [T] L_W + L_{C_1C} \\ L (\pi_K^6) &\leq L_{IB_1} + L_W \left(\frac{p_B}{\lambda_B} + \frac{p_C}{\lambda_C} \right) + L_{B_1B} + L_{BC} + L_{CG} + L_{IC_1} + L_{C_1C} \end{aligned} \quad (30)$$

Values for $\mathbb{E}_B [T]$ and $\mathbb{E}_C [T]$ were taken from Eq. (18). Here we use the fact that if $A < B + C$ is true and $A < B + D$ is true then it must also follow that $A < B + C + D$. It should be noted that this is not the maximum value for $L (\pi_K^6)$ for which policy π_K^6 is chosen but rather the upper bound estimation for that value.

2 Implementation of the System in ROS

2.1 Planning path via an Aisle

In the given map of the environment (Fig. 1a), we can see that the map has 5 labelled aisles - A, B, C, D and E. For the given system, it might be desirable to be able to control the path planning such that the planned path goes down an aisle.

To implement this system, the function *planPathToGoalViaAisle()* in *comp0037_reactive_planner_controller/src/reactive_planner_controller.py* was completed. To ensure that the path planned goes down a particular aisle, a multiple destination path planning approach was adopted such that the first destination is the mid point² of the aisle the path has go down and the second and final destination is the goal. The function *planPathToGoalViaAisle()* first checks to see if the class variable *aisleToDriveDown* is None. If this is true, it assigns the value of the function parameter **aisle** to **aisleToDriveDown** and then the coordinates of the midpoint of the aisle is found. Once, the aisle cell coordinates are obtained path planning is done by calling **planner.search** twice. First the path is planned from the starting cell (*start-CellCoords*) to the mid point of the aisle (*aisleCellCoords*). Then, a path is planned from the mid point of the aisle(*aisleCellCoords*) to the goal cell (*goalCellCoords*). Finally, the two individual paths are concatenated together and the concatenated path is plotted on search grid so that the resulting planned path can be observed.

The resulting search grid to demonstrate the performance of this function can be seen in Fig. 4 below.

²The walls dividing the aisles are disjoint in the middle this means that choosing the mid point would ensure this algorithm does not result in adding to a higher cost than is required to be able to plan a via an aisle. Hence, the mid point was chosen.




Figure 4: The demonstration of results of planning via each of the aisles.

References

Appendices